

DataStorm_Insomniacs_Final

May 30, 2024

```
[340]: import numpy as np
import pandas as pd
from sklearn import preprocessing
from sklearn.ensemble import IsolationForest
from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans
from scipy.cluster.hierarchy import dendrogram, linkage
from sklearn.decomposition import PCA

try:
    import seaborn as sns
except ImportError:
    %pip install seaborn
    import seaborn as sns

try:
    from word2number import w2n
except ImportError:
    %pip install word2number
    from word2number import w2n

import matplotlib.pyplot as plt
from scipy import stats
```

1 Loading Data

```
[341]: #Load Dataset
dataset = pd.read_csv('customer_dataset.csv')
print("----- CUSTOMER DATASET_")
↵ print("-----")
print(dataset.head())
print("-----")
```

```
----- CUSTOMER DATASET -----
  customer_id  average_monthly_spend  average_monthly_visit_frequency \
0           1                   790                        1.11
1           2                176875                        15.74
```

2	3	6812	2.11
3	4	38542	7.82
4	5	48712	7.51

	average_monthly_basket_size
0	2.84
1	33.83
2	NaN
3	10.73
4	10.04

```

c:\Users\dell\AppData\Local\Programs\Python\Python39\lib\site-
packages\IPython\core\interactiveshell.py:3553: DtypeWarning: Columns (1,2) have
mixed types.Specify dtype option on import or set low_memory=False.
exec(code_obj, self.user_global_ns, self.user_ns)

```

2 Dealing with Non Numeric Data

```

[342]: #print all non numeric columns
print("----- NON NUMERIC COLUMNS-----")
for i in dataset.columns:
    if dataset[i].dtype == 'object':
        print(i,end = " ")
print()
print("-----")
#handling non-numeric data converting them to numeric
def word_to_num(value):
    try:
        return w2n.word_to_num(value)
    except ValueError:
        if value.isdigit():
            return int(value)
        elif value.replace('.', '', 1).isdigit() and value.count('.') < 2:
            return float(value)
        return None

for column in dataset.columns:
    dataset[column] = dataset[column].apply(lambda x : word_to_num(str(x)))

```

```

----- NON NUMERIC COLUMNS-----
average_monthly_spend average_monthly_visit_frequency
-----

```

3 Preliminary Analysis of Missing Values

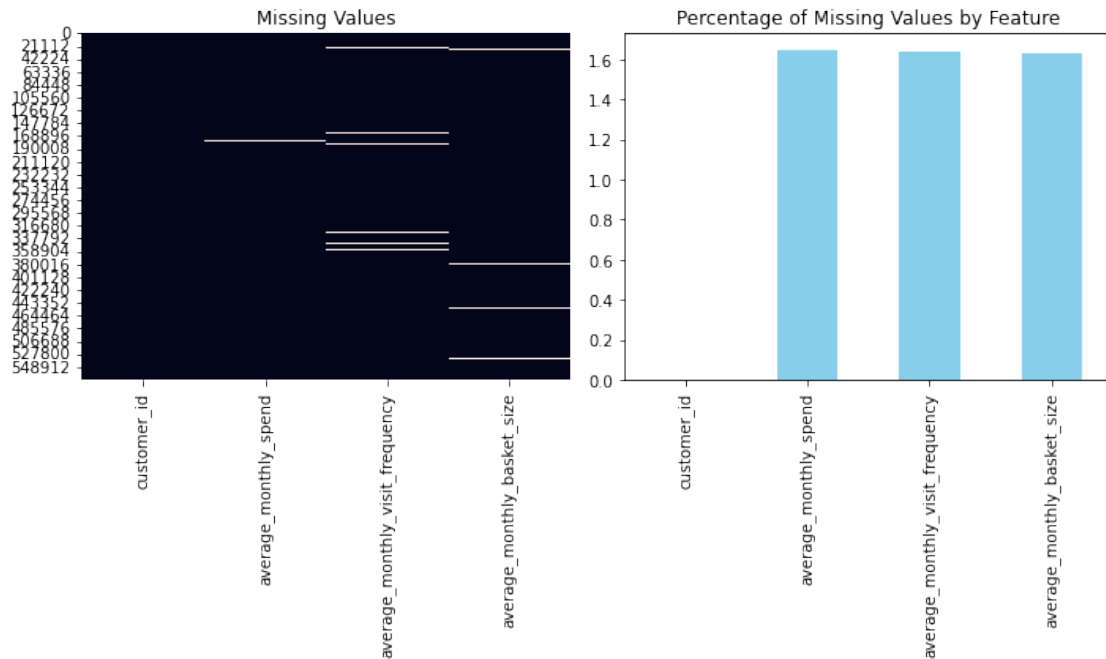
```
[343]: #missing value analysis
print("----- MISSING VALUE DETAILS_
↪-----")

#nan count
print(dataset.isna().sum())
#length of dataset
print(f"Length of dataset: {len(dataset)}")
#count all the rows with nan values
rows_with_nan = dataset.isna().any(axis=1).sum()
print(f"Rows with nan values: {rows_with_nan}")
#ratio of rows with nan values
missing_percentage = dataset.isnull().mean() * 100
print(f"Ratio of rows with nan values: {round(rows_with_nan/
↪len(dataset),2)*100}%")

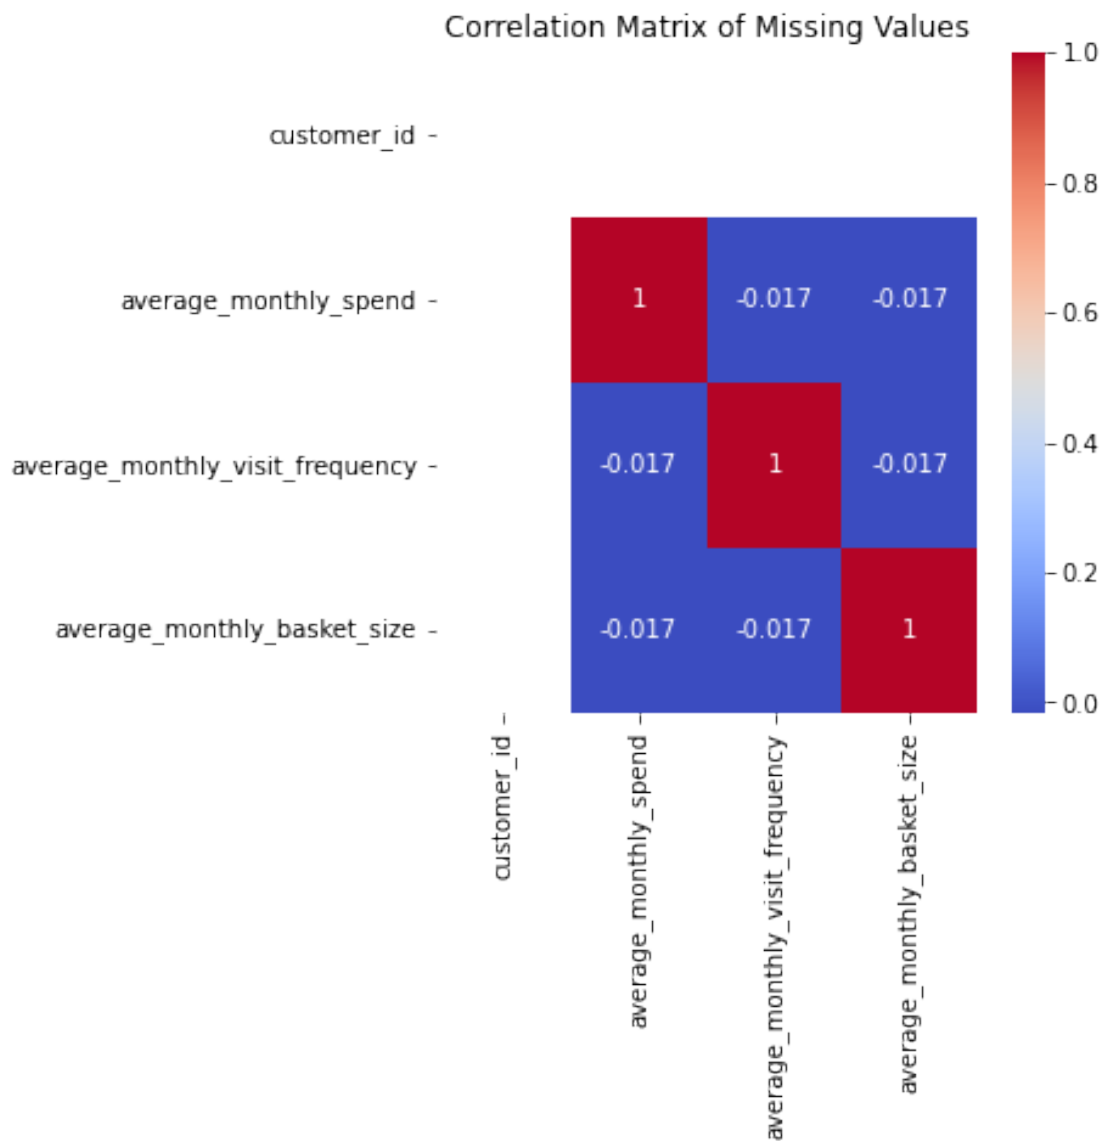
plt.figure(figsize=(5,5))
plt.subplot(1, 2, 1)
plt.title('Missing Values')
sns.heatmap(dataset.isnull(), cbar=False)

plt.subplot(1, 2, 2)
missing_percentage.plot(kind='bar', figsize=(10, 6), color='skyblue')
plt.title('Percentage of Missing Values by Feature')
plt.tight_layout()
plt.show()
print("-----
```

```
----- MISSING VALUE DETAILS
-----
customer_id                0
average_monthly_spend      9389
average_monthly_visit_frequency  9337
average_monthly_basket_size  9287
dtype: int64
Length of dataset: 570000
Rows with nan values: 28013
Ratio of rows with nan values: 5.0%
```



```
[344]: #correlation matrix to see the missing values are correlated
plt.figure(figsize=(5, 5))
corr_matrix = dataset.isnull().corr()
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm')
plt.title('Correlation Matrix of Missing Values')
plt.show()
```



```
[345]: #extract nan rows from dataset for further analysis
nan_rows = dataset[dataset.isnull().any(axis=1)]
#replace nan with 0
dataset.dropna(inplace=True)
# Handle duplicates
dataset.drop_duplicates(inplace=True)

print("----- EXTRACTED NAN ROWS -----")
print(nan_rows)
print("-----")
```

----- EXTRACTED NAN ROWS -----

	customer_id	average_monthly_spend	average_monthly_visit_frequency	\
2	3	6812.0	2.11	
7	8	12656.0	2.13	
24	25	41329.0	NaN	
52	53	171010.0	NaN	
53	54	104457.0	18.88	
...	
559914	559915	45023.0	NaN	
559945	559946	49533.0	5.87	
559959	559960	NaN	5.80	
559970	559971	95444.0	19.85	
559972	559973	13981.0	2.74	

	average_monthly_basket_size
2	NaN
7	NaN
24	10.32
52	35.03
53	NaN
...	...
559914	9.82
559945	NaN
559959	9.87
559970	NaN
559972	NaN

[28013 rows x 4 columns]

```
[346]: def plot_outliers():
    iso_forest = IsolationForest(contamination=0.05, random_state=42)
    features =
    ↪['average_monthly_spend', 'average_monthly_visit_frequency', 'average_monthly_basket_size']
    data = dataset.copy()
    iso_forest.fit(data[features])
    outliers = iso_forest.predict(data[features])
    data['outlier'] = outliers
    outliers_data = data[data['outlier'] == -1]
    sns.pairplot(data, hue='outlier', palette={1: 'blue', -1: 'red'},
    ↪markers=["o", "s"])
    plt.suptitle('Isolation Forest Outlier Detection', y=1.02)
    plt.show()
```

```
[347]: #Handle outliers
data = dataset.copy()
Q1 = dataset.quantile(0.25)
Q3 = dataset.quantile(0.75)
```

```

IQR = Q3 - Q1
outliers = ((dataset < (Q1 - 1.5 * IQR)) | (dataset > (Q3 + 1.5 * IQR))).
↳any(axis=1)
dataset = dataset[~outliers]

copy_dataset = dataset.copy()
data.describe()

```

```

[347]:
      customer_id  average_monthly_spend  average_monthly_visit_frequency \
count  531990.000000          5.319900e+05          531990.000000
mean   280059.070229          4.825899e+04           6.879288
std    161689.312262          9.173203e+04           5.323171
min         1.000000          0.000000e+00           0.000000
25%    139927.250000          1.049900e+04           2.390000
50%    280109.500000          3.195900e+04           5.600000
75%    420171.750000          8.541800e+04           9.850000
max     560000.000000          4.833000e+07          432.740000

      average_monthly_basket_size
count          531990.000000
mean              13.292845
std              12.039045
min               0.000000
25%               3.800000
50%               9.790000
75%              17.480000
max              385.040000

```

4 Data Visualization and Preliminary Clustering for Observing the Nature of Decision Boundaries

```

[348]: #scale for cluster visualization
scaler = StandardScaler()
scaled_data = scaler.fit_transform(dataset.iloc[:, 1:])
scaled_data_nan = scaler.transform(nan_rows.iloc[:, 1:])

```

```

[349]: #clustering using Kmeans
kmeans = KMeans(n_clusters=5, random_state=42)
clusters = kmeans.fit_predict(scaled_data)

```


```

c:\Users\dell\AppData\Local\Programs\Python\Python39\lib\site-
packages\sklearn\cluster\_kmeans.py:870: FutureWarning: The default value of
`n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init`
explicitly to suppress the warning
  warnings.warn(

```

```
[350]: nan_rows.fillna(0, inplace=True)
dataset['Cluster'] = clusters
fig = plt.figure(figsize=(13, 12))
ax = fig.add_subplot(121, projection='3d')
ax.scatter(dataset['average_monthly_spend'],
           dataset['average_monthly_visit_frequency'],
           dataset['average_monthly_basket_size'], c=clusters, cmap='viridis')
ax.set_xlabel('average_monthly_spend')
ax.set_ylabel('average_monthly_visit_frequency')
ax.set_zlabel('average_monthly_basket_size')
plt.title('Clusters')
ax.set_xlabel('average_monthly_spend')
ax.set_ylabel('average_monthly_visit_frequency')
ax.set_zlabel('average_monthly_basket_size')

plt.show()
```



Cuboidal shape of the clusters is proof that the Decision boundaries can be represented as a logical

combination of conditions on each parameter separately. (The decision Boundaries are aligned with the Vertical and Horizontal Planes). That is the Function for a Decision Boundary for a cluster j , $\phi_j(x_1, x_2, x_3)$ (where x_1, x_2, x_3 are the three features) can be expressed as a logical combination of three simple decision rules.

That is if,

$$\phi_j(x_1, x_2, x_3) = \begin{cases} 1 & \text{If belongs to cluster } j \\ 0 & \text{If does not belongs to cluster } j \end{cases}$$

then,

$$\phi_j(x_1, x_2, x_3) = \phi_{1,j}(x_1) \cdot \phi_{2,i}(x_2) \cdot \phi_{3,i}(x_3)$$

where \cdot is the logical and operation furthermore since there is only one such block per cluster per axis,

$$\phi_{k,j}(x_k) = \begin{cases} 1 & a_k < x_k < b_k \\ 0 & \text{otherwise} \end{cases}$$

Where a_k and b_k are to be determined Finally it can also be observed that these margins are far apart so clustering in each dimension separately is sufficient for identifying these boundaries.


Finally this corroborates well with the Problem Statement

5 Demonstrating that Rows with Missing Values Belong to One of the Clusters and Can be Fully Determined by Considering the remaining Columns

We will use the fact that if we make one of the features zero, this produces a Projection of one of the Cuboidal Clusters onto one of the planes. Therefore, we will replace the missing values with zero and plot in a 3d space

```
[351]: nan_rows.fillna(0, inplace=True)
dataset['Cluster'] = clusters
fig = plt.figure(figsize=(13, 12))
ax = fig.add_subplot(121, projection='3d')
ax.scatter(dataset['average_monthly_spend'],
           dataset['average_monthly_visit_frequency'],
           dataset['average_monthly_basket_size'], c=clusters, cmap='viridis')
ax.set_xlabel('average_monthly_spend')
ax.set_ylabel('average_monthly_visit_frequency')
ax.set_zlabel('average_monthly_basket_size')
plt.title('Clusters')
ax.scatter(nan_rows['average_monthly_spend'],
           nan_rows['average_monthly_visit_frequency'],
           nan_rows['average_monthly_basket_size'],
           c='black', s = 0.1)
ax.set_xlabel('average_monthly_spend')
ax.set_ylabel('average_monthly_visit_frequency')
ax.set_zlabel('average_monthly_basket_size')
```

```
plt.show()
```



We can easily see that all of the missing value clouds (Black data clouds) are simply projections of the already identified clusters onto one of the planes x_1x_2, x_1x_3, x_2x_3 thus proving our claim is correct. We will reinforce this claim in later sections. Finally, since all Black Data clouds are 2 Dimensional it is also clear that there are no columns with two or more missing values.

6 Demonstrating that the Distribution Roughly Follows the Rules given in the Problem Statement

```
[352]: def cluster_visualization_on_feature(feature_name:str):
        feature_dataset = copy_dataset[[feature_name]].dropna()
        n = 5 if feature_name == "average_monthly_visit_frequency" else 4
        kmeans = KMeans(n_clusters=n, random_state=42)
        clusters = kmeans.fit_predict(feature_dataset)
```

```

feature_dataset[f'{feature_name} Cluster'] = clusters
plt.figure(figsize=(5, 5))
sns.boxplot(x=f'{feature_name} Cluster', y=feature_name,
data=feature_dataset)
plt.title(f'{feature_name} Cluster')
plt.show()

cluster_visualization_on_feature('average_monthly_spend')
cluster_visualization_on_feature('average_monthly_visit_frequency')
cluster_visualization_on_feature('average_monthly_basket_size')

```


c:\Users\dell\AppData\Local\Programs\Python\Python39\lib\site-packages\sklearn\cluster_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning

```
warnings.warn(
```

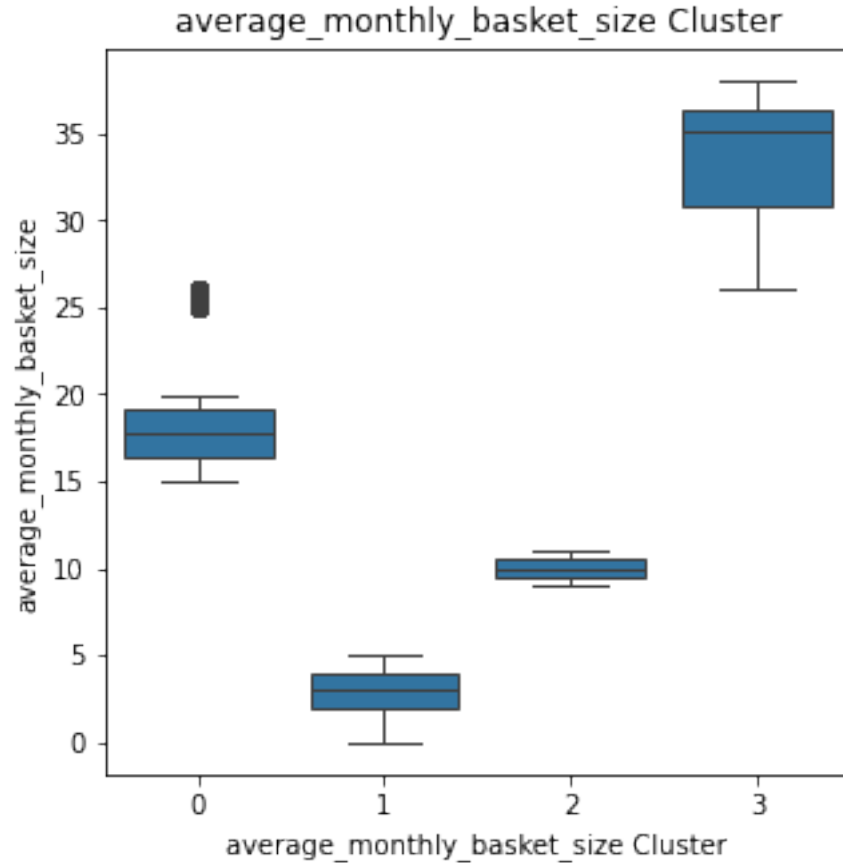


c:\Users\dell\AppData\Local\Programs\Python\Python39\lib\site-packages\sklearn\cluster_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init`

```
explicitly to suppress the warning
warnings.warn(
```



```
c:\Users\dell\AppData\Local\Programs\Python\Python39\lib\site-
packages\sklearn\cluster\_kmeans.py:870: FutureWarning: The default value of
`n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init`
explicitly to suppress the warning
warnings.warn(
```



Plotting Box and Whisker graphs to analyze distinct clusters within features. Considering Average Monthly Spend by Cluster, We can identify there are 4 distinct clusters (cluster 0 and 4 are the same). We could categorize them as, - 3 - LOW - 1 - Average - 0, 4 - High - 2 - Very High

Considering Average Monthly Visit Frequency by cluster, We can identify that there are 5 distinct clusters. We can categorize them as,

- 3 - Low
- 1 - Average Lower
- 4 - Average Higher
- 2 - High
- 1 - Very High

Note that we can see two well separated clusters in the average range instead of the one as given in the problem statement. It is possible that these two clusters have no bearing on the problem. But we will later show that these have a very high importance for the classification task. Considering Average Monthly Basket Size by cluster. Although we can see 5 clusters in the plot we can see that the two upper clusters are overlapping in terms of the inter quartile ranges. Therefore will only consider four clusters. We can categorize them as ,

- 3 - Low
- 1 - Average

- 0 - High
- 2,4 - Very High

The Table is mostly correct however there are some modifications need to be done which we will further verify and update in a later section

7 Analyzing Each Feature Set Independantly and Identifying the Percentiles (Ranges in our Case)

Since in the previous section we demonstrated that it is sufficient to analyze each feature independantly we wil analyze them, - first visually through histograms - second using

```
[353]: print("-----Data Analysis-----")

data = pd.read_csv('customer_dataset.csv')
data.head()
```

```
-----Data Analysis-----

c:\Users\dell\AppData\Local\Programs\Python\Python39\lib\site-
packages\IPython\core\interactiveshell.py:3553: DtypeWarning: Columns (1,2) have
mixed types.Specify dtype option on import or set low_memory=False.
  exec(code_obj, self.user_global_ns, self.user_ns)
```

```
[353]:  customer_id average_monthly_spend average_monthly_visit_frequency \
0          1              790              1.11
1          2          176875          15.74
2          3           6812           2.11
3          4          38542           7.82
4          5          48712           7.51

      average_monthly_basket_size
0              2.84
1             33.83
2              NaN
3             10.73
4             10.04
```

```
[354]: # Extracting columns from the dataframe
col_names = data.columns.tolist()
print(col_names)
```

```
['customer_id', 'average_monthly_spend', 'average_monthly_visit_frequency',
'average_monthly_basket_size']
```

```
[355]: # Converting words to numerical values
def convert_word_to_number(amount):
    try:
```

```

        return w2n.word_to_num(amount)
    except ValueError:
        return amount

for col in col_names:
    data[col] = data[col].apply(convert_word_to_number)
    data[col] = pd.to_numeric(data[col], errors='coerce')
    data = data.dropna(subset=[col])

```

```

[356]: def plot_barCharts(column_name, x_axis_label, interval_width : int, min_limit : int, max_limit : int):
    max_value = int(data[column_name].max()) + interval_width

    bins = range(min_limit, max_limit, interval_width)
    labels = [f'{i}-{i+interval_width}' for i in bins[:-1]]

    data['range'] = pd.cut(data[column_name], bins=bins, labels=labels, right=False)
    customer_count = data['range'].value_counts().sort_index()

    plt.figure(figsize=(12, 6))
    customer_count.plot(kind='bar')
    plt.xlabel(x_axis_label)
    plt.ylabel('Number of Customers')
    plt.title(f'Number of Customers in Each {x_axis_label} Range')
    plt.xticks(rotation=45)
    plt.tight_layout()
    plt.show()


    print(customer_count)

```

```

[357]: plot_barCharts('average_monthly_spend', 'Average Monthly Spend', 5000, 0, 200001)

```




0-5000	64687
5000-10000	64284
10000-15000	64552
15000-20000	0
20000-25000	32134
25000-30000	32692
30000-35000	32145
35000-40000	32006
40000-45000	32371
45000-50000	32357
50000-55000	0
55000-60000	0
60000-65000	0
65000-70000	0
70000-75000	0
75000-80000	0
80000-85000	17794
85000-90000	17918
90000-95000	17680
95000-100000	17782
100000-105000	17605
105000-110000	17601
110000-115000	0
115000-120000	0
120000-125000	0
125000-130000	0
130000-135000	0
135000-140000	0
140000-145000	6057

145000-150000	5982
150000-155000	6146
155000-160000	6188
160000-165000	6173
165000-170000	6061
170000-175000	5856
175000-180000	5918
180000-185000	0
185000-190000	0
190000-195000	0
195000-200000	0

Name: range, dtype: int64

```
[358]: plot_barCharts('average_monthly_visit_frequency', 'Average Monthly Visit_
↪Frequency', 1, 0, 50)
```

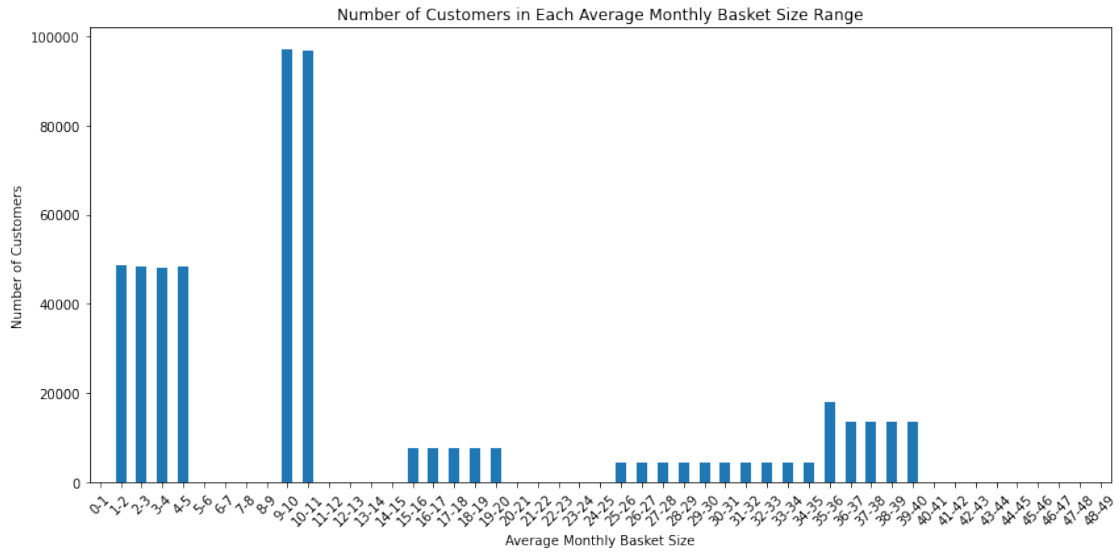


0-1	2
1-2	96944
2-3	96575
3-4	1
4-5	48285
5-6	48346
6-7	48647
7-8	48423
8-9	0
9-10	22447
10-11	22718
11-12	22519
12-13	0

13-14	0
14-15	24151
15-16	24235
16-17	0
17-18	0
18-19	12928
19-20	25766
20-21	0
21-22	0
22-23	1
23-24	0
24-25	0
25-26	0
26-27	0
27-28	0
28-29	0
29-30	0
30-31	0
31-32	0
32-33	0
33-34	0
34-35	0
35-36	0
36-37	0
37-38	0
38-39	0
39-40	0
40-41	0
41-42	0
42-43	0
43-44	0
44-45	0
45-46	0
46-47	0
47-48	0
48-49	0

Name: range, dtype: int64

```
[359]: plot_barCharts('average_monthly_basket_size', 'Average Monthly Basket Size', 1,
↪0, 50)
```



0-1	3
1-2	48568
2-3	48458
3-4	48130
4-5	48366
5-6	0
6-7	0
7-8	0
8-9	0
9-10	97051
10-11	96656
11-12	0
12-13	0
13-14	0
14-15	0
15-16	7824
16-17	7632
17-18	7771
18-19	7709
19-20	7758
20-21	0
21-22	0
22-23	0
23-24	0
24-25	0
25-26	4400
26-27	4505
27-28	4346
28-29	4374

29-30	4480
30-31	4440
31-32	4356
32-33	4408
33-34	4343
34-35	4352
35-36	17985
36-37	13562
37-38	13471
38-39	13460
39-40	13587
40-41	0
41-42	0
42-43	0
43-44	0
44-45	0
45-46	0
46-47	0
47-48	0
48-49	0

Name: range, dtype: int64

8 Further Proving that rows with Missing Values Follow the Same Distribution as the other ones

```
[360]: # Now we'll try to analyze data by removing only the N/A entries in the
        ↳ respective column
```

```
dataset = pd.read_csv('customer_dataset.csv')
```

```
c:\Users\dell\AppData\Local\Programs\Python\Python39\lib\site-
packages\IPython\core\interactiveshell.py:3553: DtypeWarning: Columns (1,2) have
mixed types.Specify dtype option on import or set low_memory=False.
  exec(code_obj, self.user_global_ns, self.user_ns)
```

```
[361]: col_names = dataset.columns.tolist()
        print(col_names)
```

```
['customer_id', 'average_monthly_spend', 'average_monthly_visit_frequency',
'average_monthly_basket_size']
```

```
[362]: for col in col_names:
        dataset[col] = dataset[col].apply(convert_word_to_number)
        dataset[col] = pd.to_numeric(dataset[col], errors='coerce')
```

```
[363]: def plot_barCharts_2(column_name, x_axis_label, interval_width : int, min_limit,
↳: int, max_limit : int):
    data = dataset.dropna(subset=[column_name])
    max_value = int(data[column_name].max()) + interval_width

    bins = range(min_limit, max_limit, interval_width)
    labels = [f'{i}-{i+interval_width}' for i in bins[:-1]]

    data['range'] = pd.cut(data[column_name], bins=bins, labels=labels,
↳right=False)
    customer_count = data['range'].value_counts().sort_index()

    plt.figure(figsize=(12, 6))
    customer_count.plot(kind='bar')
    plt.xlabel(x_axis_label)
    plt.ylabel('Number of Customers')
    plt.title(f'Number of Customers in Each {x_axis_label} Range')
    plt.xticks(rotation=45)
    plt.tight_layout()
    plt.show()

    print(customer_count)
```

```
[364]: plot_barCharts_2('average_monthly_spend', 'Average Monthly Spend', 5000, 0,
↳200001)
```

C:\Users\dell\AppData\Local\Temp\ipykernel_26116\4100806091.py:8:


SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
data['range'] = pd.cut(data[column_name], bins=bins, labels=labels,
right=False)
```



0-5000	66892
5000-10000	66494
10000-15000	66722
15000-20000	0
20000-25000	33268
25000-30000	33795
30000-35000	33218
35000-40000	33105
40000-45000	33501
45000-50000	33424
50000-55000	0
55000-60000	0
60000-65000	0
65000-70000	0
70000-75000	0
75000-80000	0
80000-85000	18421
85000-90000	18560
90000-95000	18285
95000-100000	18413
100000-105000	18226
105000-110000	18248
110000-115000	0
115000-120000	0
120000-125000	0
125000-130000	0
130000-135000	0
135000-140000	0
140000-145000	6254


145000-150000	6205
150000-155000	6355
155000-160000	6380
160000-165000	6394
165000-170000	6270
170000-175000	6058
175000-180000	6115
180000-185000	0
185000-190000	0
190000-195000	0
195000-200000	0

Name: range, dtype: int64

```
[365]: plot_barCharts_2('average_monthly_visit_frequency', 'Average Monthly Visit_
↪Frequency', 1, 0, 50)
```

C:\Users\dell\AppData\Local\Temp\ipykernel_26116\4100806091.py:8:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
data['range'] = pd.cut(data[column_name], bins=bins, labels=labels, right=False)



0-1	2
1-2	100263
2-3	99976

3-4	1
4-5	49959
5-6	49982
6-7	50306
7-8	50022
8-9	0
9-10	23255
10-11	23506
11-12	23352
12-13	0
13-14	0
14-15	25007
15-16	25036
16-17	0
17-18	0
18-19	13356
19-20	26636
20-21	0
21-22	0
22-23	1
23-24	0
24-25	0
25-26	0
26-27	0
27-28	0
28-29	0
29-30	0
30-31	0
31-32	0
32-33	0
33-34	0
34-35	0
35-36	0
36-37	0
37-38	0
38-39	0
39-40	0
40-41	0
41-42	0
42-43	0
43-44	0
44-45	0
45-46	0
46-47	0
47-48	0
48-49	0

Name: range, dtype: int64


```
[366]: plot_barCharts_2('average_monthly_basket_size', 'Average Monthly Basket Size', 1, 0, 50)
```

C:\Users\dell\AppData\Local\Temp\ipykernel_26116\4100806091.py:8:


SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
data['range'] = pd.cut(data[column_name], bins=bins, labels=labels, right=False)
```



0-1	3
1-2	50256
2-3	50105
3-4	49792
4-5	50098
5-6	0
6-7	0
7-8	0
8-9	0
9-10	100359
10-11	99922
11-12	0
12-13	0
13-14	0
14-15	0
15-16	8092

16-17	7899
17-18	8084
18-19	7970
19-20	8032
20-21	0
21-22	0
22-23	0
23-24	0
24-25	0
25-26	4553
26-27	4639
27-28	4499
28-29	4526
29-30	4619
30-31	4598
31-32	4519
32-33	4598
33-34	4519
34-35	4493
35-36	18557
36-37	14019
37-38	13959
38-39	13938
39-40	14061
40-41	0
41-42	0
42-43	0
43-44	0
44-45	0
45-46	0
46-47	0
47-48	0
48-49	0

Name: range, dtype: int64

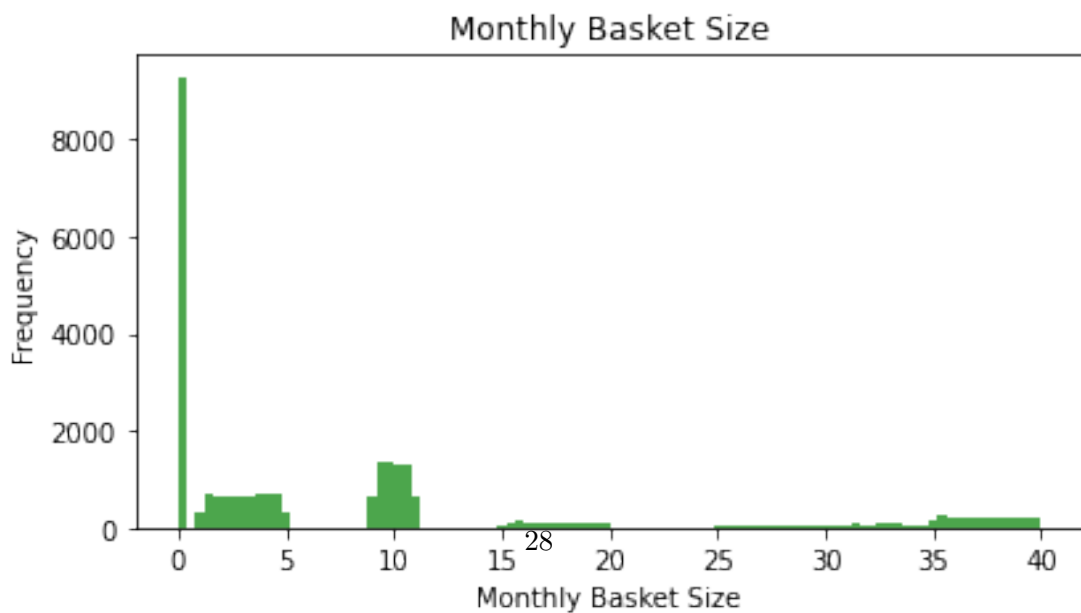
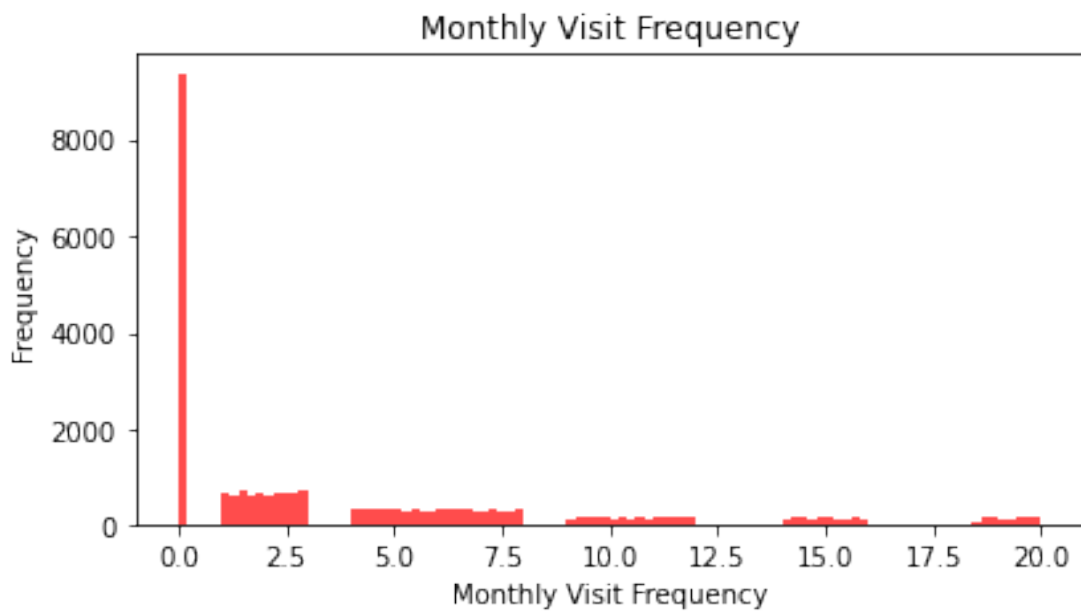

```
[367]: plt.figure(figsize=(6, 10))
plt.subplots_adjust(hspace=1)
plt.subplot(3, 1, 1)
plt.hist(nan_rows['average_monthly_spend'], bins=100, color='blue', alpha=0.7)
plt.title('Monthly Spend')
plt.xlabel('Monthly Spend')
plt.ylabel('Frequency')

plt.subplot(3, 1, 2)
plt.hist(nan_rows['average_monthly_visit_frequency'], bins=100, color='red',
        alpha=0.7)
```

```
plt.title('Monthly Visit Frequency')
plt.xlabel('Monthly Visit Frequency')
plt.ylabel('Frequency')

plt.subplot(3, 1, 3)
plt.hist(nan_rows['average_monthly_basket_size'], bins=100, color='green',
        ↪alpha=0.7)
plt.title('Monthly Basket Size')
plt.xlabel('Monthly Basket Size')
plt.ylabel('Frequency')

plt.tight_layout()
plt.show()
```



9 Feature Selection

9.1 Loading Data

```
[368]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from word2number import w2n
dataset_og = pd.read_csv("customer_dataset.csv")
CUS_ID = 'customer_id'
AVG_SPEND = 'average_monthly_spend'
AVG_VISIT = 'average_monthly_visit_frequency'
AVG_BASKET_SIZE = 'average_monthly_basket_size'
#handling non-numeric data converting them to numeric
def word_to_num(value):
    try:
        return w2n.word_to_num(value)
    except ValueError:
        if value.isdigit():
            return int(value)
        elif value.replace('.', '', 1).isdigit() and value.count('.') < 2:
            return float(value)
        return None

for column in dataset_og.columns:
    dataset_og[column] = dataset_og[column].apply(lambda x :
↪word_to_num(str(x)))

dataset = dataset_og.drop_duplicates().dropna()
dataset.head()
```

```
c:\Users\dell\AppData\Local\Programs\Python\Python39\lib\site-
packages\IPython\core\interactiveshell.py:3553: DtypeWarning: Columns (1,2) have
mixed types.Specify dtype option on import or set low_memory=False.
    exec(code_obj, self.user_global_ns, self.user_ns)
```

```
[368]:
```

	customer_id	average_monthly_spend	average_monthly_visit_frequency	\
0	1	790.0	1.11	
1	2	176875.0	15.74	
3	4	38542.0	7.82	
4	5	48712.0	7.51	
5	6	172333.0	14.16	

average_monthly_basket_size

0	2.84
1	33.83
3	10.73
4	10.04
5	35.06

9.2 Defining Cluster Thresholds Identified Preivously

```
[369]: class Basket_Sizes:
    LOW = (0,7)
    AVERAGE = (7,13)
    HIGH = (13,22)
    VERY_HIGH = (22,50)
    ranges = [LOW,AVERAGE,HIGH,VERY_HIGH]
    ranges_good = [LOW,AVERAGE,HIGH,VERY_HIGH]

    @classmethod
    def filter_set(cls,dataset,range):
        assert(range in cls.ranges)
        return dataset[(dataset[AVG_BASKET_SIZE] >= range[0]) &
↪(dataset[AVG_BASKET_SIZE] < range[1])]

    @classmethod
    def get_mask(cls,dataset,range):
        assert(range in cls.ranges)
        return ((dataset[AVG_BASKET_SIZE] >= range[0]) &
↪(dataset[AVG_BASKET_SIZE] < range[1]))

class Monthly_Spend:
    LOW = (0,17500)
    AVERAGE = (17500,65000)
    HIGH = (65000,125000)
    VERY_HIGH = (125000,200000)
    ranges = [LOW,AVERAGE,HIGH,VERY_HIGH]
    ranges_good = [LOW,AVERAGE,HIGH,VERY_HIGH]

    @classmethod
    def filter_set(cls,dataset,range):
        assert(range in cls.ranges)
        return dataset[(dataset[AVG_SPEND] >= range[0]) & (dataset[AVG_SPEND] <
↪range[1])]

    @classmethod
    def get_mask(cls,dataset,range):
        assert(range in cls.ranges)
```

```

        return ((dataset[AVG_SPEND] >= range[0]) & (dataset[AVG_SPEND] <
↪range[1]))

class Average_Visit:
    LOW = (0,3.7)
    AVERAGE_LOWER = (3.7,8.5)
    AVERAGE_HIGHER = (8.5,13)
    AVERAGE = (3.7,13)
    HIGH = (13,17)
    VERY_HIGH = (17,21)
    ranges = [LOW,AVERAGE_LOWER,AVERAGE_HIGHER,HIGH,VERY_HIGH,AVERAGE]
    ranges_good = [LOW,AVERAGE_LOWER,AVERAGE_HIGHER,HIGH,VERY_HIGH]

    @classmethod
    def filter_set(cls,dataset,range):
        assert(range in cls.ranges)
        return dataset[(dataset[AVG_VISIT] >= range[0]) & (dataset[AVG_VISIT] <
↪range[1])]
    @classmethod
    def get_mask(cls,dataset,range):
        assert(range in cls.ranges)
        return ((dataset[AVG_VISIT] >= range[0]) & (dataset[AVG_VISIT] <
↪range[1]))

```

9.3 Checking Limiting and Borderline Cases in Thresholds Manually

While the Clusters boarder that we have identified programatically (and verified by visual inspection) have good clear boarders there are a few data points that is better if closely and manually verified to ensure it has been given the correct group

```
[370]: data[(data[AVG_VISIT]>3) & (data[AVG_VISIT] < 4)]
```

```
[370]:
```

	customer_id	average_monthly_spend	average_monthly_visit_frequency \
569999	244801	4225.0	3.4

	average_monthly_basket_size	range
569999	1.92	1-2

```
[371]: data[(data[AVG_VISIT]>3) & (data[AVG_VISIT] < 4)]
```

```
[371]:
```

	customer_id	average_monthly_spend	average_monthly_visit_frequency \
569999	244801	4225.0	3.4

	average_monthly_basket_size	range
569999	1.92	1-2

9.4 Filtering Using the Exact Rules Given in the Problem Statement

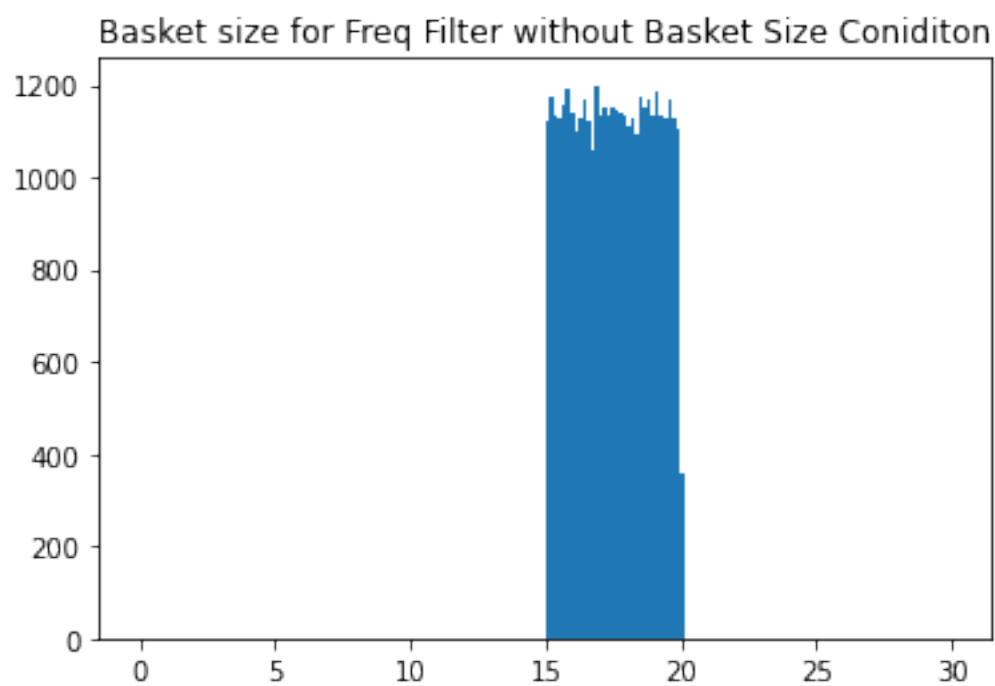
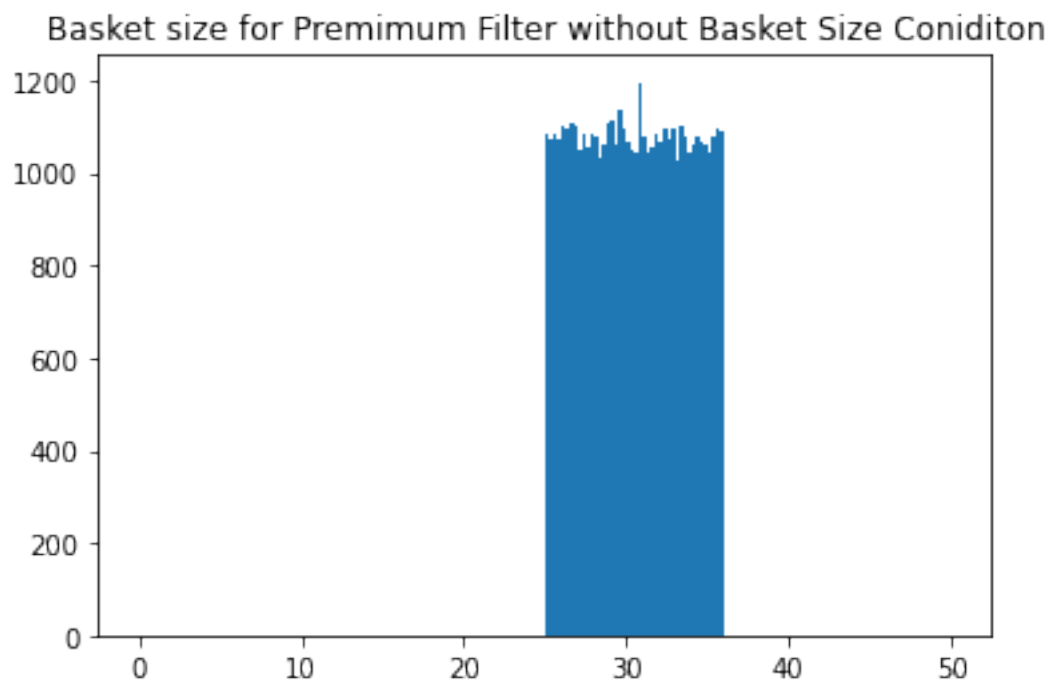
```
[372]: low_level_mask = Monthly_Spend.get_mask(dataset,Monthly_Spend.LOW) &
        ↳Average_Visit.get_mask(dataset,Average_Visit.LOW) & Basket_Sizes.
        ↳get_mask(dataset,Basket_Sizes.LOW)
medium_level_mask = Monthly_Spend.get_mask(dataset,Monthly_Spend.AVERAGE) &
        ↳Average_Visit.get_mask(dataset,Average_Visit.AVERAGE) & Basket_Sizes.
        ↳get_mask(dataset,Basket_Sizes.AVERAGE)
high_frequent_mask = Monthly_Spend.get_mask(dataset,Monthly_Spend.HIGH) &
        ↳Average_Visit.get_mask(dataset,Average_Visit.VERY_HIGH) & Basket_Sizes.
        ↳get_mask(dataset,Basket_Sizes.AVERAGE)
high_loyal_mask = Monthly_Spend.get_mask(dataset,Monthly_Spend.HIGH) &
        ↳Average_Visit.get_mask(dataset,Average_Visit.AVERAGE) & Basket_Sizes.
        ↳get_mask(dataset,Basket_Sizes.VERY_HIGH)
high_premium_mask = Monthly_Spend.get_mask(dataset,Monthly_Spend.VERY_HIGH) &
        ↳Average_Visit.get_mask(dataset,Average_Visit.HIGH) & Basket_Sizes.
        ↳get_mask(dataset,Basket_Sizes.HIGH)
print(low_level_mask.sum())
print(medium_level_mask.sum())
print(high_frequent_mask.sum())
print(high_loyal_mask.sum())
print(high_premium_mask.sum())
```

```
189981
190117
0
66407
0
```

From this we have verified the calim that we have made before that the basket value for the “High-end-frequent” and “High-end-premium” needs to be changed from the values given in the table in the problem statement. This can be verified by taking the datasets that do not fit the rules and plotting the distribution of the datasets that do not follow the rules and apply the rules for these two datasets except for the Basket size dataset and observing the distribution

```
[373]: okay_masks = [low_level_mask,medium_level_mask,high_loyal_mask]
other_data = dataset[~(okay_masks[0] | okay_masks[1] | okay_masks[2])]

plt.title("Basket size for Premium Filter without Basket Size Coniditon")
plt.hist(Monthly_Spend.filter_set(other_data,Monthly_Spend.
        ↳VERY_HIGH) [AVG_BASKET_SIZE],bins=200,range=(0,50))
plt.show()
plt.title("Basket size for Freq Filter without Basket Size Coniditon")
plt.hist(Monthly_Spend.filter_set(other_data,Monthly_Spend.
        ↳HIGH) [AVG_BASKET_SIZE],bins=200,range=(0,30))
plt.show()
```

[374] :

```

high_frequent_mask = Monthly_Spend.get_mask(dataset,Monthly_Spend.HIGH) &
    ↳Average_Visit.get_mask(dataset,Average_Visit.VERY_HIGH) & Basket_Sizes.
    ↳get_mask(dataset,Basket_Sizes.HIGH)
high_premium_mask = Monthly_Spend.get_mask(dataset,Monthly_Spend.VERY_HIGH) &
    ↳Average_Visit.get_mask(dataset,Average_Visit.HIGH) & Basket_Sizes.
    ↳get_mask(dataset,Basket_Sizes.VERY_HIGH)
dataset[high_frequent_mask].describe()

```

```

[374]:
count    customer_id  average_monthly_spend  average_monthly_visit_frequency \
count    37975.000000          37975.000000          37975.000000
mean     281239.489849          94997.415036          19.247284
std      161976.207837           8660.220172           0.434948
min        74.000000          80000.000000          18.500000
25%     140928.000000          87549.500000          18.870000
50%     281125.000000          95004.000000          19.250000
75%     421880.000000         102523.000000          19.630000
max     559984.000000         109999.000000          19.990000

        average_monthly_basket_size
count                37975.000000
mean                  17.492247
std                   1.443141
min                   15.000000
25%                   16.240000
50%                   17.500000
75%                   18.750000
max                   19.990000

```

9.5 Proving additional clusters Identified in the Average Frequency correspond to actual features

Although the problem statement only identifies one average cluster for Frequency parameter from our preliminary investigations we have shown that two separate clusters are observable and furthermore that they actually can be effectively used to identify customer type. Now we will prove this claim. By applying the other two rules for a dataset with outliers removed and plotting the histograms


```



[375]: data = dataset
def remove_outliers(data, column):
    Q1 = data[column].quantile(0.25)
    Q3 = data[column].quantile(0.75)
    IQR = Q3 - Q1
    data = data[~((data[column] < (Q1 - 1.5 * IQR)) | (data[column] > (Q3 + 1.5
    ↳* IQR)))]
    return data
data = remove_outliers(data,AVG_SPEND)
data = remove_outliers(data,AVG_VISIT)

```

```
data = remove_outliers(data,AVG_BASKET_SIZE)
```

```
[376]: data1 = Monthly_Spend.filter_set(Basket_Sizes.filter_set(data,Basket_Sizes.  
    ↳LOW),Monthly_Spend.LOW)  
plt.hist(data1[AVG_VISIT],bins=100)  
plt.show()  
data2 = Monthly_Spend.filter_set(Basket_Sizes.filter_set(data,Basket_Sizes.  
    ↳AVERAGE),Monthly_Spend.AVERAGE)  
plt.hist(data2[AVG_VISIT],bins=100)  
plt.show()  
data3 = Monthly_Spend.filter_set(Basket_Sizes.filter_set(data,Basket_Sizes.  
    ↳VERY_HIGH),Monthly_Spend.HIGH)  
plt.hist(data3[AVG_VISIT],bins=100)  
plt.show()
```





In fact keeping these two as separate clusters is highly beneficial especially since this alone can be used to identify between Medium and High-End Loyal Customers even if the two other parameters (Average Spending and Average Basket size) are missing.

This leads us to conclude that the table should be updated as | Customer Category | Monthly Spend

| Monthly Visit Frequency | Monthly Basket Size | |-----|-----|-----|
 ---|-----| | High end - Premium | Very High(3) | High(3) | High Very High(3) | | High
 end - Loyal | High(2) | Average (Average_Higher)(2) | Very High(3) | | High end - Frequent | High(2)
 | Very High(4) | Average High(2) | | Medium Level | Average(1) | Average (Average_Lower)(1) |
 Average(1) | | Low Level | Low(0) | Low(0) | Low(0) |

10 Analyzing Nature of Outliers and their Following of Rules

Now we will consider the dataset without the outliers to see if they obey the patterns observed

```
[377]: other_data = dataset[~(low_level_mask | medium_level_mask | high_loyal_mask |
    ↪high_frequent_mask | high_premium_mask )]
other_data.describe()
```

```
[377]:
```

	customer_id	average_monthly_spend	average_monthly_visit_frequency \
count	18.000000	1.800000e+01	18.000000
mean	354434.666667	4.934703e+06	46.058333
std	159311.270356	1.293920e+07	100.784055
min	252.000000	0.000000e+00	1.250000
25%	322879.500000	9.543150e+04	11.930000
50%	325078.000000	1.489220e+05	15.310000
75%	515194.750000	1.612676e+06	19.630000
max	558101.000000	4.833000e+07	432.740000

	average_monthly_basket_size
count	18.000000
mean	69.655000
std	110.472051
min	2.160000
25%	12.427500
50%	29.285000
75%	37.470000
max	385.040000

Most of the values that do not follow the rules are due to outliers so we will get rid of the outliers and see if there are columns which follow the rules

```
[378]: spend_anomaly = other_data[AVG_SPEND]>Monthly_Spend.VERY_HIGH[1]
basket_anomaly = other_data[AVG_BASKET_SIZE]>Basket_Sizes.VERY_HIGH[1]
visit_anomaly = other_data[AVG_VISIT]>Average_Visit.VERY_HIGH[1]

other_data[~(spend_anomaly | basket_anomaly | visit_anomaly)]
```

```
[378]:
```

	customer_id	average_monthly_spend	average_monthly_visit_frequency \
325525	325526	0.0	9.36
515037	515038	0.0	15.35

	average_monthly_basket_size
--	-----------------------------

325525	36.42
515037	30.02

From simple observation alone we can see that these two simply correspond to an entering error in the average monthly spend and actually belong to two clear categories if we consider the other two features

11 Feature Engineering

Now we will give a numeric label to each of the ranges for each of the parameters (known as the score) Outliers and Nan values are marked as known as (-1). (These can actually later be estimated atleast the range of them by considering the cluster that it belongs to)

```
[379]: classification_data = dataset_og.copy()
VISIT_SCORE = 'VISIT_SCORE'
BASKET_SCORE = 'BASKET_SCORE'
SPEND_SCORE = 'SPEND_SCORE'

classification_data[VISIT_SCORE] = -1
classification_data[BASKET_SCORE] = -1
classification_data[SPEND_SCORE] = -1

for i,rr in enumerate(Monthly_Spend.ranges_good):
    mask = Monthly_Spend.get_mask(classification_data,rr)
    classification_data.loc[mask,SPEND_SCORE] = i
    #do the same for other two

for i,rr in enumerate(Basket_Sizes.ranges_good):
    mask = Basket_Sizes.get_mask(classification_data,rr)
    classification_data.loc[mask,BASKET_SCORE] = i

for i,rr in enumerate(Average_Visit.ranges_good):
    mask = Average_Visit.get_mask(classification_data,rr)
    classification_data.loc[mask,VISIT_SCORE] = i

classification_data
```

```
[379]:
```

	customer_id	average_monthly_spend	average_monthly_visit_frequency \
0	1	790.0	1.11
1	2	176875.0	15.74
2	3	6812.0	2.11
3	4	38542.0	7.82
4	5	48712.0	7.51
...
569995	341740	27428.0	4.91

569996	215276	141241.0	15.47
569997	11515	9183.0	2.09
569998	205260	35541.0	4.22
569999	244801	4225.0	3.40

	average_monthly_basket_size	VISIT_SCORE	BASKET_SCORE	SPEND_SCORE
0	2.84	0	0	0
1	33.83	3	3	3
2	NaN	0	-1	0
3	10.73	1	1	1
4	10.04	1	1	1
...
569995	9.42	1	1	1
569996	31.13	3	3	3
569997	4.61	0	0	0
569998	9.65	1	1	1
569999	1.92	0	0	0

[570000 rows x 7 columns]

```
[380]: print(classification_data[VISIT_SCORE].value_counts())
print(classification_data[SPEND_SCORE].value_counts())
print(classification_data[BASKET_SCORE].value_counts())
```

```
1    200269
0    200242
2     70113
3     50043
4     39992
-1     9341
Name: VISIT_SCORE, dtype: int64
1     200311
0     200108
2     110153
3       50031
-1       9397
Name: SPEND_SCORE, dtype: int64
1     200281
0     200254
3     120097
2       40077
-1        9291
Name: BASKET_SCORE, dtype: int64
```

Demonstrate that **There are no data points that have all features unknown**

```
[381]:
```

```
no_data_mask = (classification_data[VISIT_SCORE]==-1) &
↳(classification_data[SPEND_SCORE] == -1) &
↳(classification_data[BASKET_SCORE] == -1)
classification_data[no_data_mask]
```

[381]: Empty DataFrame
Columns: [customer_id, average_monthly_spend, average_monthly_visit_frequency, average_monthly_basket_size, VISIT_SCORE, BASKET_SCORE, SPEND_SCORE]
Index: []

Our Final Algorithm of Classification is very simple.

We have already identified the rules for each cluster by considering the Problem Statement Table, The K-means Cluster Whisker plots. Fact that the decision boundaries can be written as logistic boundaries using and combination of decision inequalities for components independently (demonstrated using the cuboid shape of the cluster). The Nan Values also follow the same Law since they can clearly be seen to be projections of these cubes. Therefore we can simply collapse all features by assigning a number as for each range. Then we find the similarity to a certain Cluster rules using the following formula.

Let the similarity between a customer and a Cluster be defined by $S(\mathbf{x}_i, \mathbf{c}_j)$ Where \mathbf{x}_i is the feature vector defined for the i^{th} customer, and \mathbf{c}_j for the j^{th} cluster. A feature vector is defined as [Average Spend Cluster Index, Average Basket Size Cluster Index, Average Visit Frequency Index]

Now,

$$S(x_i, c_j) = \sum_{k=0}^3 (x_i[k] == c_j[k])$$

where k stands for the k^{th} feature

That is the count of the features where the cluster and customer has equal feature Indexes

Now the Cluster of the i^{th} customer is given as,

$$\text{Cluster of the } i^{th} \text{ customer} = \arg \max_{j=0}^5 S(x_i, c_j)$$

11.1 Calculating Cluster Similarities

```
[382]: score_given_data = classification_data.copy()

PREMIUM = 'PREMIUM'
LOYAL = 'LOYAL'
FREQ = 'FREQ'
MED = 'MED'
LOW = 'LOW'

PREMIUM_VEC = {SPEND_SCORE:3,VISIT_SCORE:3,BASKET_SCORE:3}
LOYAL_VEC = {SPEND_SCORE:2,VISIT_SCORE:2,BASKET_SCORE:3}
```



```

FREQ_VEC = {SPEND_SCORE:2,VISIT_SCORE:4,BASKET_SCORE:2}
MED_VEC = {SPEND_SCORE:1,VISIT_SCORE:1,BASKET_SCORE:1}
LOW_VEC = {SPEND_SCORE:0,VISIT_SCORE:0,BASKET_SCORE:0}

Types = [PREMIUM,LOYAL,FREQ,MED,LOW]
Vecs = [PREMIUM_VEC,LOYAL_VEC,FREQ_VEC,MED_VEC,LOW_VEC]
for ct,vec in zip(Types,Vecs):
    scores = (score_given_data[SPEND_SCORE]== vec[SPEND_SCORE]).astype(int)
    ↪(score_given_data[BASKET_SCORE]==vec[BASKET_SCORE]).astype(int)
    ↪(score_given_data[VISIT_SCORE]==vec[VISIT_SCORE]).astype(int)
    print(ct)
    score_given_data[ct] = scores

#get the column name that has the maximum value or values out of Types

MAX_CUSTOMER_TYPE_SCORE = 'MAX_SCORE'
score_given_data[MAX_CUSTOMER_TYPE_SCORE] =
    ↪score_given_data[[PREMIUM,LOYAL,FREQ,MED,LOW]].max(axis=1)

```

```

PREMIUM
LOYAL
FREQ
MED
LOW

```

11.2 Calculaing the Cluster Type

```

[383]: final_result = score_given_data.copy()
CUSTOMER_TYPE = 'CUSTOMER_TYPE'

final_result[CUSTOMER_TYPE] = final_result[[PREMIUM,LOYAL,FREQ,MED,LOW]].
    ↪idxmax(axis=1)
final_result[CUSTOMER_TYPE].value_counts()

```

```

[383]: MED          203581
LOW          203542
LOYAL         71275
PREMIUM       50887
FREQ          40715
Name: CUSTOMER_TYPE, dtype: int64

```

```

[384]: final_result[[CUS_ID,CUSTOMER_TYPE]]

```

```

[384]:      customer_id CUSTOMER_TYPE
0              1          LOW
1              2        PREMIUM

```

2	3	LOW
3	4	MED
4	5	MED
...
569995	341740	MED
569996	215276	PREMIUM
569997	11515	LOW
569998	205260	MED
569999	244801	LOW

[570000 rows x 2 columns]

12 Saving the Data

```
[385]: saving_data = final_result.copy()[[CUS_ID,CUSTOMER_TYPE]]
assert(len(dataset_og)==len(saving_data))#making sure all the data is there
saving_data[CUSTOMER_TYPE] = saving_data[CUSTOMER_TYPE].map({PREMIUM: 'High end_
↳ Premium',LOYAL: 'High end - Loyal',FREQ: 'High end - Frequent',MED: 'Medium_
↳ level',LOW: 'Low level'})
saving_data.to_csv("final_results.csv",index=False)
saving_data.head()
```

```
[385]:  customer_id      CUSTOMER_TYPE
0         1          Low level
1         2  High end - Premium
2         3          Low level
3         4      Medium level
4         5      Medium level
```

13 Checking the Results

```
[386]: validating_set = final_result.copy()
```

13.1 Proving that there are no Customers who have more than one compatible cluster

```
[387]: duplicates = (score_given_data[Types[0]] ==_
↳ score_given_data[MAX_CUSTOMER_TYPE_SCORE]).astype(int)
for type in Types[1:]:
    duplicates = (duplicates +(score_given_data[type] ==_
↳ score_given_data[MAX_CUSTOMER_TYPE_SCORE]).astype(int) )
duplicates.value_counts()
```

```
[387]: 1    570000
      dtype: int64
```

Since all results have only one Max score there is no ambiguity

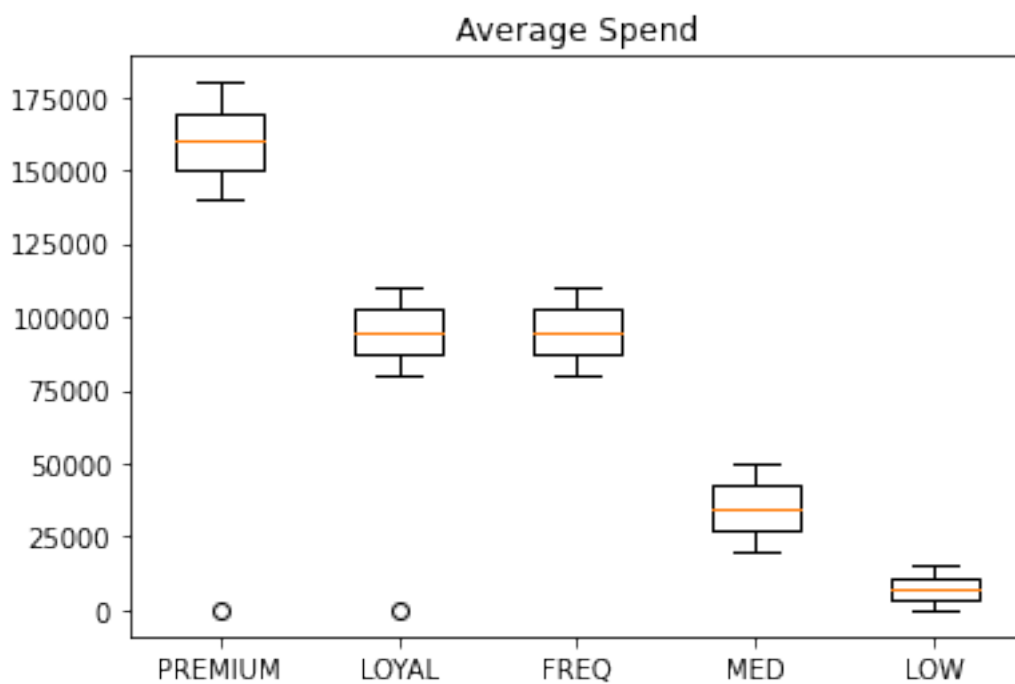
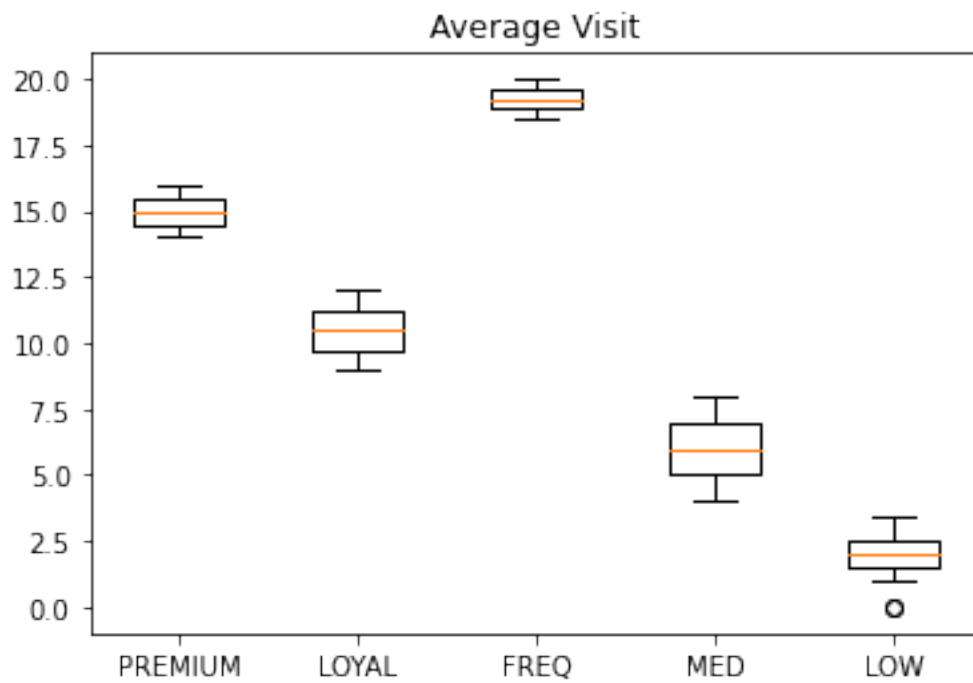
```
[ ]:
```

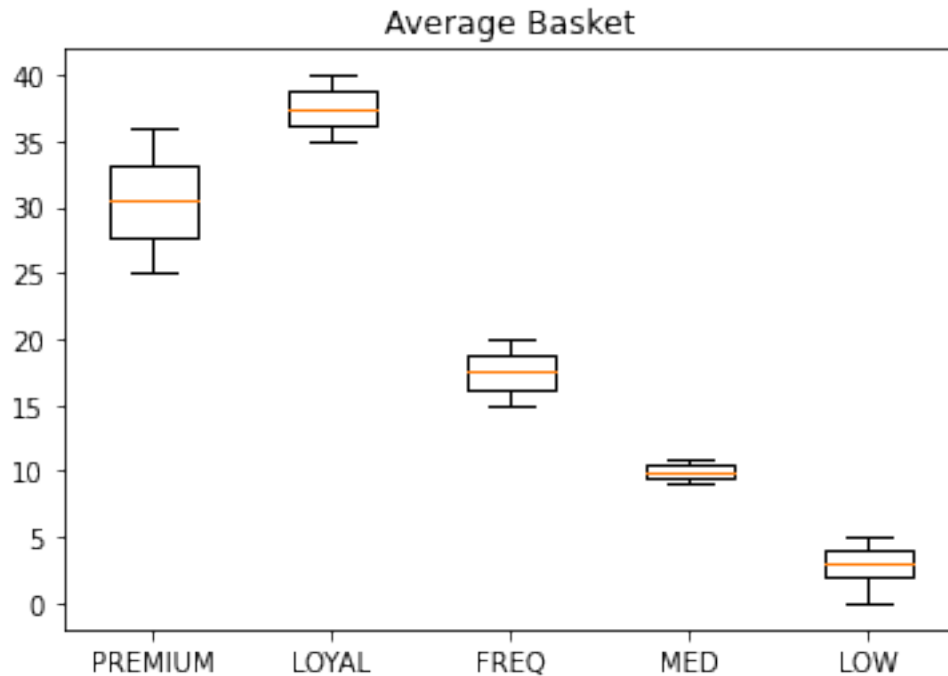
13.2 Checking that the Distributions of Feature Values with the Final Classification Match with Expectation

```
[388]: plt.title("Average Visit")
plt.boxplot(
    [validating_set[validating_set[CUSTOMER_TYPE]==type][AVG_VISIT] .
    ↪dropna()[validating_set[validating_set[CUSTOMER_TYPE]==type][AVG_VISIT] .
    ↪dropna().between(0,Average_Visit.VERY_HIGH[1])] for type in Types],
    labels=Types
)
plt.show()

plt.title("Average Spend")
plt.boxplot(
    [validating_set[validating_set[CUSTOMER_TYPE]==type][AVG_SPEND] .
    ↪dropna()[validating_set[validating_set[CUSTOMER_TYPE]==type][AVG_SPEND] .
    ↪dropna().between(0,Monthly_Spend.VERY_HIGH[1])] for type in Types],
    labels=Types
)
plt.show()

plt.title("Average Basket")
plt.boxplot(
    [validating_set[validating_set[CUSTOMER_TYPE]==type][AVG_BASKET_SIZE] .
    ↪dropna()[validating_set[validating_set[CUSTOMER_TYPE]==type][AVG_BASKET_SIZE] .
    ↪dropna().between(0,Basket_Sizes.VERY_HIGH[1])] for type in Types],
    labels=Types
)
plt.show()
```





14 Checking the Nature of the customers who had only similarity of 2

```
[389]: score_is_2 = validating_set[validating_set[MAX_CUSTOMER_TYPE_SCORE]==2]
score_is_2
```

```
[389]:
```

	customer_id	average_monthly_spend	average_monthly_visit_frequency	\
2	3	6812.0	2.11	
7	8	12656.0	2.13	
24	25	41329.0	NaN	
52	53	171010.0	NaN	
53	54	104457.0	18.88	
...	
559914	559915	45023.0	NaN	
559945	559946	49533.0	5.87	
559959	559960	NaN	5.80	
559970	559971	95444.0	19.85	
559972	559973	13981.0	2.74	

	average_monthly_basket_size	VISIT_SCORE	BASKET_SCORE	SPEND_SCORE	\
2	NaN	0	-1	0	
7	NaN	0	-1	0	
24	10.32	-1	1	1	

52	35.03	-1	3	3
53	NaN	4	-1	2
...
559914	9.82	-1	1	1
559945	NaN	1	-1	1
559959	9.87	1	1	-1
559970	NaN	4	-1	2
559972	NaN	0	-1	0

	PREMIUM	LOYAL	FREQ	MED	LOW	MAX_SCORE	CUSTOMER_TYPE
2	0	0	0	0	2	2	LOW
7	0	0	0	0	2	2	LOW
24	0	0	0	2	0	2	MED
52	2	1	0	0	0	2	PREMIUM
53	0	1	2	0	0	2	FREQ
...
559914	0	0	0	2	0	2	MED
559945	0	0	0	2	0	2	MED
559959	0	0	0	2	0	2	MED
559970	0	1	2	0	0	2	FREQ
559972	0	0	0	0	2	2	LOW

[28031 rows x 14 columns]

```
[390]: score_is_2.dropna().describe()
```

```
[390]:
```

	customer_id	average_monthly_spend	average_monthly_visit_frequency \
count	18.000000	1.800000e+01	18.000000
mean	354434.666667	4.934703e+06	46.058333
std	159311.270356	1.293920e+07	100.784055
min	252.000000	0.000000e+00	1.250000
25%	322879.500000	9.543150e+04	11.930000
50%	325078.000000	1.489220e+05	15.310000
75%	515194.750000	1.612676e+06	19.630000
max	558101.000000	4.833000e+07	432.740000

	average_monthly_basket_size	VISIT_SCORE	BASKET_SCORE	SPEND_SCORE \
count	18.000000	18.000000	18.000000	18.000000
mean	69.655000	1.666667	1.277778	0.388889
std	110.472051	1.847096	1.673515	1.539247
min	2.160000	-1.000000	-1.000000	-1.000000
25%	12.427500	0.000000	0.000000	-1.000000
50%	29.285000	2.500000	1.500000	0.000000
75%	37.470000	3.000000	3.000000	2.000000
max	385.040000	4.000000	3.000000	3.000000

PREMIUM	LOYAL	FREQ	MED	LOW	MAX_SCORE
---------	-------	------	-----	-----	-----------

count	18.000000	18.000000	18.000000	18.000000	18.000000	18.0
mean	0.888889	0.722222	0.444444	0.222222	0.444444	2.0
std	0.963382	0.751904	0.783823	0.646762	0.783823	0.0
min	0.000000	0.000000	0.000000	0.000000	0.000000	2.0
25%	0.000000	0.000000	0.000000	0.000000	0.000000	2.0
50%	0.500000	1.000000	0.000000	0.000000	0.000000	2.0
75%	2.000000	1.000000	0.750000	0.000000	0.750000	2.0
max	2.000000	2.000000	2.000000	2.000000	2.000000	2.0

We see that the ones that have given a similairyt score of 2 is the ones that are outliers the, two entires which had zero as spending be mistake (This is the above 18 rows), and entires with NA. Since we already know that the ones with NA only have one NA it is reasonable that all have a similarity score of 2. Furthermore it shows that even there there is no ambiguity

Thus our classification is Justified